

STAT 455 Project

2025-11-29

Read/inspect data

Source: https://www.eia.gov/dnav/ng/ng_enr_drill_s1_m.htm

```
library(ggplot2)
library(forecast)

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

library(zoo)

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
## 
##   as.Date, as.Date.numeric

library(acp)

## Loading required package: tseries

## Loading required package: quantmod

## Loading required package: xts

## Loading required package: TTR

library(tscount)
library(ggplot2)
library(dplyr)

##
## ##### Warning from 'xts' package #####
## #
## # The dplyr lag() function breaks how base R's lag() function is supposed to #
## # work, which breaks lag(my_xts). Calls to lag(my_xts) that you type or      #
## # source() into this session won't work correctly.                            #
```

```

## #
## # Use stats::lag() to make sure you're not using dplyr::lag(), or you can add #
## # conflictRules('dplyr', exclude = 'lag') to your .Rprofile to stop #
## # dplyr from breaking base R's lag() function. #
## #
## # Code in packages is not affected. It's protected by R's namespace mechanism #
## # Set 'options(xts.warn_dplyr_breaks_lag = FALSE)' to suppress this warning. #
## #
## #####
## Attaching package: 'dplyr'

## The following objects are masked from 'package:xts':
##   first, last

## The following objects are masked from 'package:stats':
##   filter, lag

## The following objects are masked from 'package:base':
##   intersect, setdiff, setequal, union


data <- read.csv('rigs.csv')

data$date <- as.Date(paste0("01-", data$date), format = "%d-%b-%y")

data %>%
  ggplot(aes(date, count)) +
  geom_line() +
  theme_minimal() +
  labs(x = 'Date',
       y = 'Count',
       title = 'US Oil Rig Count by Date')

```

US Oil Rig Count by Date



Convert data to time series

```
data$date <- as.yearmon(data$date, "%b-%y")

z <- ts(data$count,
         start = c(as.numeric(format(data$date[1], "%Y")),
                    as.numeric(format(data$date[1], "%m"))),
         frequency = 12)
```


Specifying test set

```
# use a 5-year test window
test_years <- 5
test_horizon <- test_years * 12


n <- length(z)
start_test <- n - test_horizon
h <- 1 # 1-step forecast
```

ACF and PACF Plots

```
acf(as.numeric(z)[1:start_test], main = "ACF")
```



```
pacf(as.numeric(z)[1:start_test], main = "PACF")
```



Running models and evaluating

```
# Function to return metrics
evaluate_metrics <- function(actual, forecasted) {
  acc <- accuracy(forecasted, actual)
  data.frame(
    ME    = acc["Test set", "ME"],
    RMSE = acc["Test set", "RMSE"],
    MAE   = acc["Test set", "MAE"],
    MPE   = acc["Test set", "MPE"],
    MAPE  = acc["Test set", "MAPE"]
  )
}

# Function to plot forecast vs actual
evaluate_plot <- function(actual, forecasted, dates, title = "Forecast vs Actual") {
  df <- data.frame(
    date = rep(dates, 2),
    value = c(actual, forecasted),
    type = rep(c("Actual", "Forecasted"), each = length(actual))
  )

  ggplot(df, aes(x = date, y = value, color = type)) +
    geom_line()
}
```

```

geom_line() +
theme_minimal() +
labs(
  x = "Date",
  y = "Count",
  title = title
) +
theme(legend.title = element_blank())
}

# Function to plot errors
error_plot <- function(actual, forecasted, dates, title = "Forecast Error") {
  df <- data.frame(
    date = dates,
    error = actual - forecasted
  )

  ggplot(df, aes(x = date, y = error)) +
  geom_line(color = "red") +
  geom_hline(yintercept = 0, linetype = "dashed") +
  theme_minimal() +
  labs(
    x = "Date",
    y = "Error (Actual - Forecast)",
    title = title
  )
}

# Setup
test_years <- 5
test_horizon <- test_years * 12
n <- length(z)
start_test <- n - test_horizon
h <- 1
dates <- index(z)

results_list <- list()      # store metrics
forecast_plots <- list()    # store forecast vs actual plots
error_plots <- list()      # store error plots

# ARIMA models
arima_models <- list(
  AR1      = c(1,1,0),
  MA1      = c(0,1,1),
  ARMA11  = c(1,1,1)
)

for (name in names(arima_models)) {
  pred <- rep(NA, n)

  for (i in start_test:(n - h)) {
    train <- window(z, end = time(z)[i])
    fit <- tryCatch(Arima(train, order = arima_models[[name]]), error = function(e) NULL)
  }
}

```

```

    if (!is.null(fit)) pred[i + h] <- forecast(fit, h = h)$mean[h]
}


idx <- start_test + h
actual <- z[idx:n]
forecasted <- pred[idx:n]
valid <- !is.na(forecasted)

# Metrics
metrics <- evaluate_metrics(actual[valid], forecasted[valid])
metrics <- cbind(Model = name, metrics)
results_list[[name]] <- metrics


# Forecast vs actual plot
p_forecast <- evaluate_plot(actual[valid], forecasted[valid], dates[idx:n][valid],
                             title = paste("ARIMA -", name, "Forecast"))
print(p_forecast)
forecast_plots[[name]] <- p_forecast

# Error plot
p_error <- error_plot(actual[valid], forecasted[valid], dates[idx:n][valid],
                       title = paste("ARIMA -", name, "Forecast Error"))
print(p_error)
error_plots[[name]] <- p_error
}


```




ARIMA – AR1 Forecast Error




ARIMA – MA1 Forecast




ARIMA – MA1 Forecast Error



ARIMA – ARMA11 Forecast



ARIMA – ARMA11 Forecast Error



```
# TSGLM models
z_scaled <- z / 100

tsglm_models <- list(
  AR1      = list(past_obs = 1),
  MA1      = list(past_mean = 1),
  ARMA11   = list(past_obs = 1, past_mean = 1)
)

for (name in names(tsglm_models)) {
  pred <- rep(NA, n)
  max_lag <- max(unlist(tsglm_models[[name]]))
  start <- max(start_test, max_lag + 5)

  for (i in start:(n - h)) {
    train <- z_scaled[1:i]
    fit <- suppressWarnings(
      tryCatch(
        tsglm(train, model = tsglm_models[[name]], link = "identity", distr = "poisson",
              init.method = "marginal", init.drop = FALSE),
        error = function(e) NULL
      )
    )
    lambda_next <- if (!is.null(fit)) predict(fit, n.ahead = h)$pred[h] else mean(train)
    pred[i + h] <- lambda_next * 100
  }
}
```

```


idx <- start + h
actual <- z[idx:n]
forecasted <- pred[idx:n]
valid <- !is.na(forecasted)

# Metrics
metrics <- evaluate_metrics(actual[valid], forecasted[valid])
metrics <- cbind(Model = paste0("TSGLM_", name), metrics)
results_list[[paste0("TSGLM_", name)]] <- metrics


# Forecast vs actual plot
p_forecast <- evaluate_plot(actual[valid], forecasted[valid], dates[idx:n][valid],
                             title = paste("TSGLM -", name, "Forecast"))
print(p_forecast)
forecast_plots[[paste0("TSGLM_", name)]] <- p_forecast

# Error plot
p_error <- error_plot(actual[valid], forecasted[valid], dates[idx:n][valid],
                       title = paste("TSGLM -", name, "Forecast Error"))
print(p_error)
error_plots[[paste0("TSGLM_", name)]] <- p_error
}


```




TSGLM – AR1 Forecast Error




TSGLM – MA1 Forecast




TSGLM – MA1 Forecast Error



TSGLM – ARMA11 Forecast



TSGLM – ARMA11 Forecast Error



```
# Combine results
comparison_table <- do.call(rbind, results_list)
rownames(comparison_table) <- NULL

comparison_table
```

##	Model	ME	RMSE	MAE	MPE	MAPE
## 1	AR1	5.053356	18.86039	14.93825	1.1935105	2.745920
## 2	MA1	4.986523	18.62340	14.74655	1.1775409	2.710276
## 3	ARMA11	4.897235	18.43444	14.58590	1.1597454	2.679870
## 4	TSGLM_AR1	1.014541	16.08024	12.79012	0.4412349	2.299766
## 5	TSGLM_MA1	-769.216588	778.51376	769.21659	-137.8471038	137.847104
## 6	TSGLM_ARMA11	-0.803470	15.63198	12.53037	0.0741234	2.189756